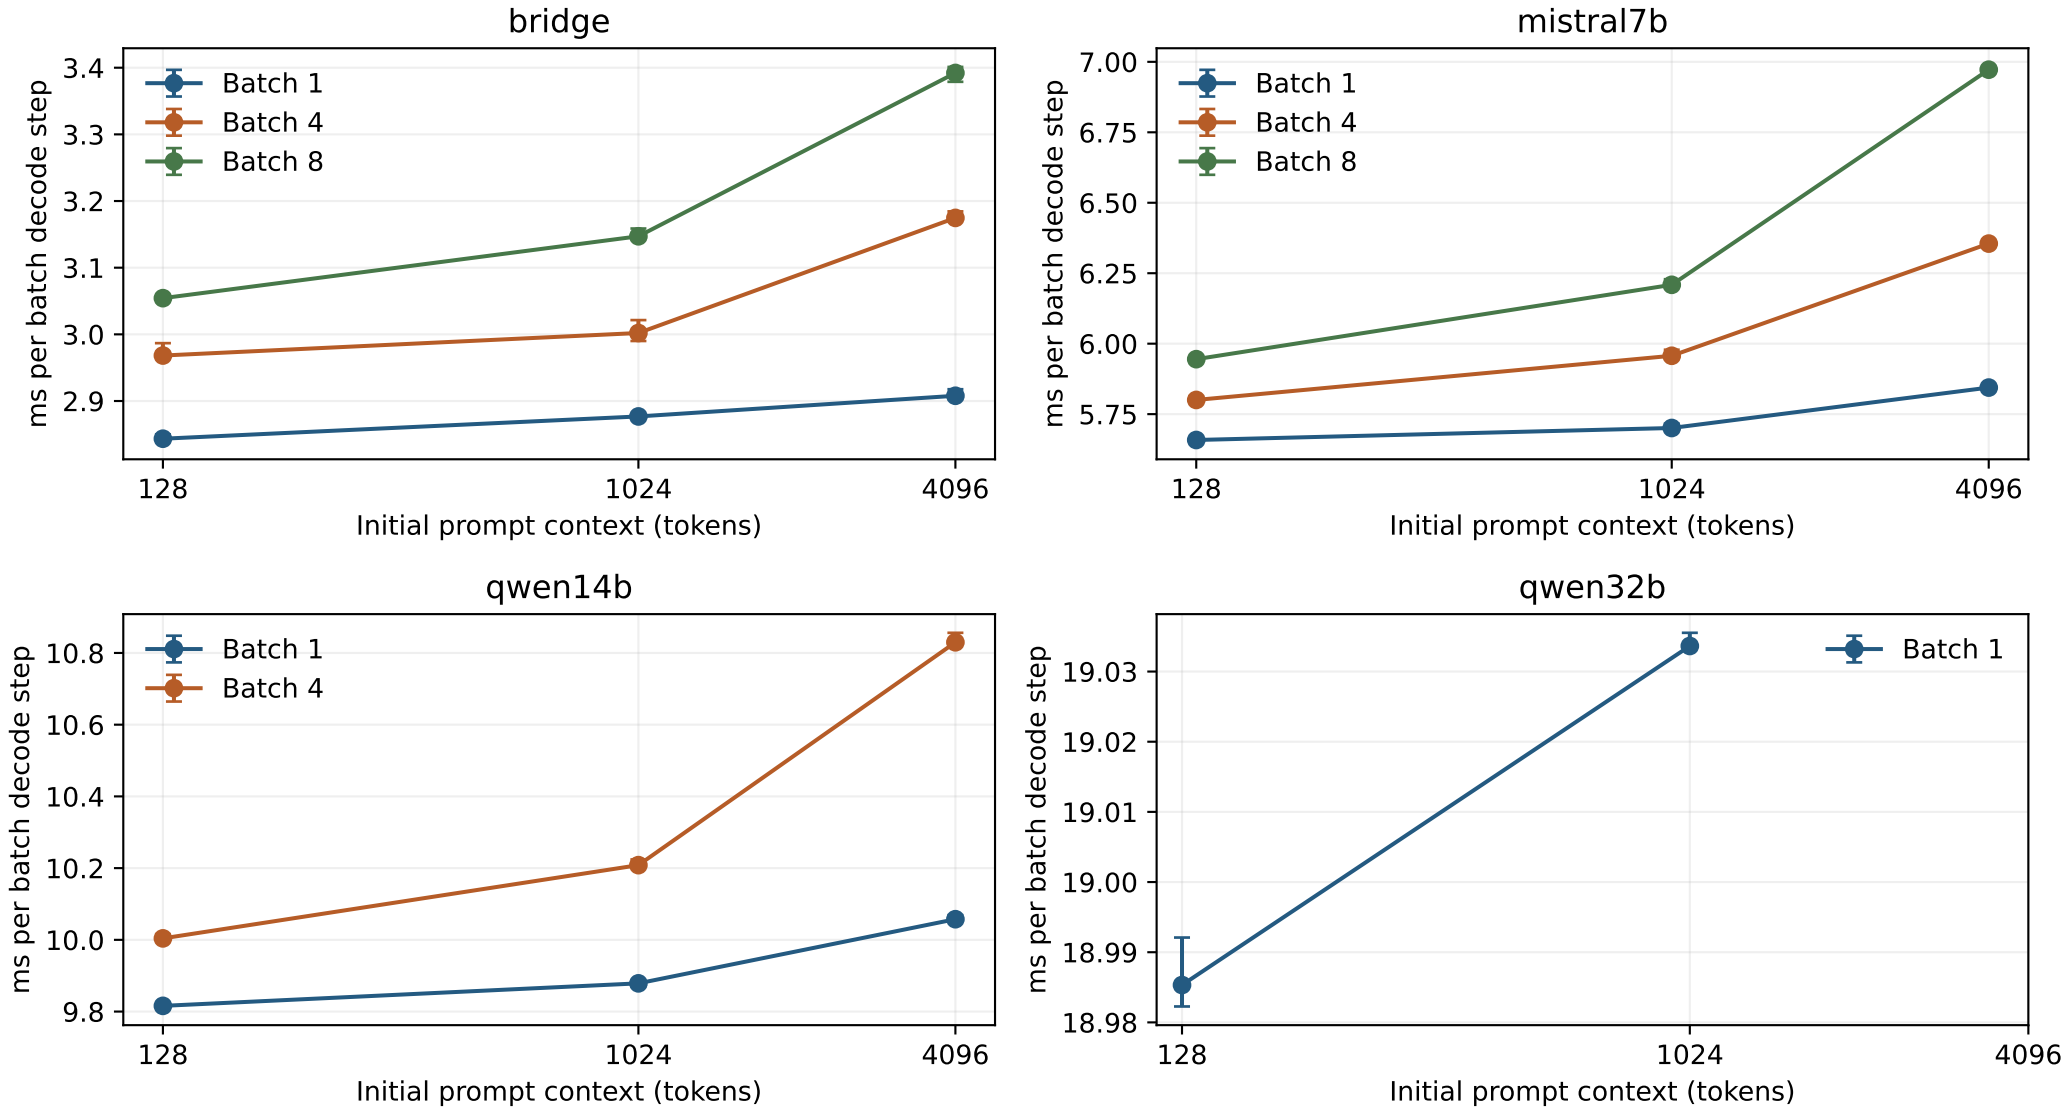
H200 controlled BF16 graph decode · measured-real-GPU
bridge
mistral7b
qwen14b
qwen32b
Batch 1
Batch 4
Batch 8
ms per batch decode step
Initial prompt context (tokens)
128
1024
4096
3.4
3.3
3.2
3.1
3.0
2.9
7.00
6.75
6.50
6.25
6.00
5.75
10.8
10.6
10.4
10.2
10.0
9.8
19.03
19.02
19.01
19.00
18.99
18.98
Three warmups and three trials; error bars show trial min–max. 64 generated tokens, 63 timed steps; prefill excluded.
Only qualified cases are plotted. Different models are different workloads.